

Lucy Li

Natural language processing (NLP), computational social science, cultural analytics, AI fairness

✉ lucy3_li@berkeley.edu | 🏠 lucy3.github.io

Education

University of California, Berkeley

Berkeley, CA

PhD Information Science

Aug 2019 - present

- Advisor: David Bamman
- Committee members: Isaac Bleaman, Niloufar Salehi, Dan Jurafsky
- Berkeley AI Research (BAIR), Stanford SPARQ Research Affiliate

Stanford University

Stanford, CA

BS Symbolic Systems, MS Computer Science

Sept 2014 - June 2019

- Coterminous degrees, w/ a concentration in language and depth in artificial intelligence.

Awards, Fellowships, & Grants

Rising Star in Data Science , University of Chicago & University of San Diego	2023
Rising Star in EECS , Rising Stars Academic Career Workshop	2023
Meta Research PhD Fellowship Finalist , Meta	2023
AI2 Outstanding Intern of the Year Award , Allen Institute for Artificial Intelligence	2022
Human-Centered Artificial Intelligence Seed Grant , Stanford HAI (PI: Patricia Bromley)	2021
Best Paper Award , Conference on Educational Data Science	2020
Graduate Research Fellowship , National Science Foundation	2019
K. Jon Barwise Award for Distinguished Contributions , Stanford Symbolic Systems	2018
Undergraduate Advising & Research (UAR) Small Grant , \$1500, Stanford University	2018
Grants for Education and Research , \$1145, Stanford Symbolic Systems	2017
Phi Beta Kappa , Stanford University	2017

Papers

*indicates equal contribution.

Working Papers & Preprints

Li Lucy, Suchin Gururangan, Luca Soldaini, Emma Strubell, David Bamman, Lauren Klein, Jesse Dodge. AboutMe: Using Self-Descriptions in Webpages to Document the Effects of English Pretraining Data Filters. *Arxiv*, 2024.

Luca Soldaini, Rodney Kinney, Akshita Bhagia, Dustin Schwenk, David Atkinson, Russell Authur, Ben Bogin, Khyathi Chandu, Jennifer Dumas, Yanai Elazar, Valentin Hofmann, Ananya Harsh Jha, Sachin Kumar, **Li Lucy**, Xinxin Lyu, Nathan Lambert, Ian Magnusson, Jacob Morrison, Niklas Muennighoff, Aakanksha Naik, Crystal Nam, Matthew E. Peters, Abhilasha Ravichander, Kyle Richardson, Zejiang Shen, Emma Strubell, Nishant Subramani, Oyvind Tafjord, Pete Walsh, Luke Zettlemoyer, Noah A. Smith, Hannaneh Hajishirzi, Iz Beltagy, Dirk Groeneveld, Jesse Dodge, Kyle Lo. Dolma: an Open Corpus of Three Trillion Tokens for Language Model Pretraining Research. *Arxiv*, 2024.

Minju Choi*, **Li Lucy***. "Othering" through War: Depiction of Asians/Asian Americans in U.S. History Textbooks from California and Texas. Revise and resubmit, 2023.

Journals & Conferences

Li Lucy, Jesse Dodge, David Bamman, Katherine A. Keith. Words as Gatekeepers: Measuring Discipline-specific Terms and Meanings in Scholarly Publications. *Findings of the Association of Computational Linguistics (ACL)*, 2023.

Li Lucy, Divya Tadimeti, David Bamman. Discovering Differences in the Representation of People using Contextualized Semantic Axes. *Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2022

Li Lucy, David Bamman. Characterizing English variation across social media communities with BERT. *Transactions of the Association of Computational Linguistics (TACL)*, 2021.

Li Lucy*, Dora Demszky*, Patricia Bromley, Dan Jurafsky. Content Analysis of Textbooks via Natural Language Processing: Findings on Gender, Race, and Ethnicity in Texas U.S. History Textbooks. *AERA Open*, 2020. [Best paper at American Educational Research Association (AERA) Educational Data Science Conference.]

Workshops & Non-Archival Conferences

Li Lucy, Su Lin Blodgett, Milad Shokouhi, Hanna Wallach, Alexandra Olteanu. “One-size-fits-all”? Observations and Expectations of NLG Systems across Identity-related Language Features. *New Directions in Analyzing Text as Data*, 2023. [non-archival]

Li Lucy, David Bamman. Gender and Representation Bias in GPT-3 Generated Stories. *Workshop on Narrative Understanding (WNU) at the North American Association for Computational Linguistics (NAACL)*, 2021.

Emma Lurie, **Li Lucy**, Masha Belyi, Sofia Dewar, Daniel Rincón, John Baldwin, Rajvardhan Oak. Investigating Causal Effects of Instructions in Crowdsourced Claim Matching. *Computation + Journalism Symposium (C+J)*, 2020. [non-archival]

Li Lucy, Julia Mendelsohn. Using sentiment induction to understand variation in gendered online communities. *Society for Computation in Linguistics (SCiL)*, 2019.

Li Lucy, Jon Gauthier. Are distributional representations ready for the real world? Evaluating word vectors for grounded perceptual meaning. *Language Grounding for Robotics (RoboNLP) Workshop at the Association for Computational Linguistics (ACL)*, 2017.

For Non-Research Audiences

Maria Antoniak, **Li Lucy**, Maarten Sap, Luca Soldaini. Using Large Language Models With Care. 2023. [Link.]

David Jurgens, **Li Lucy**. A Look inside the Pedagogy of Natural Language Processing. 2018. [Link.]

Presentations

AboutMe: Using Self-Descriptions in Webpages to Document the Effects of English Pretraining Data Filters

March 2024. Computing and Society Seminar, Stanford University.

Feb 2024. Computational Linguistics group, University of Toronto.

Measuring Depictions and Expressions of Social Groups with NLP

Jan 2024. Jantina Tammes School of Digital Society, Technology and AI. Groningen University, Netherlands.

May 2023. Stanford NLP Seminar, Stanford University.

Reader Response to Characters of Color in Literature Taught in U.S. Schools

Jan 2024. Center for Interdisciplinary Research. Bielefeld University, Germany.

Words as Gatekeepers: Measuring Discipline-specific Terms and Meanings in Scholarly Publications

Nov 2023. New Directions in Analyzing Text as Data (TADA). Amherst, MA.

Large-Scale Language Patterns Across Social Groups

Oct 2023. USC Natural Language Group, University of Southern California.

“One-size-fits-all”? Observations and Expectations of NLG Systems across Identity-related Language Features.

Aug 2023. AI, Ethics and Effects in Engineering and Research (Aether), Microsoft.

Context-Dependent Depictions of People Across Three Domains.

March 2023. “NLP for Social Science: From Language Models to Social Structures,” Columbia University.

What big data and big models bring to the table.

February 2023. “Roundtable Series: Learning How to Play with the Machines,” University of California, Berkeley.

Social NLP.

April 2022. Guest lecture, “Natural Language Processing,” University of California, Berkeley.

Characterizing English variation across social media communities with BERT.

Oct 2022. Guest lecture, “Practical Approaches to Data Science with Text,” Emory University.

June 2021. Guest lecture, “Computational Text Analysis,” Barnard College.

Nov 2021. Guest lecture, “Practical Approaches to Data Science with Text,” Emory University.

Content Analysis of Textbooks via Natural Language Processing.

Sept 2022. McGill Narrative and Society Conference, Montreal.

Oct 2021. 103rd Anniversary of the School of Information, Berkeley.

Feb 2021. Guest lecture, “Doing Digital History,” Stanford.

Feb 2021. Stanford Human-Computer Interaction Lunch Seminar.

May 2021. Guest lecture, “Using Data to Describe the World,” Stanford.

May 2020. Guest lecture, “Using Data to Describe the World,” Stanford.
Oct 2019. New Directions in Analyzing Text as Data (TADA). Stanford, CA.

Experience

Allen Institute for Artificial Intelligence

Research Intern

- Mentors: Jesse Dodge (June 2023 - Dec 2023), Kyle Lo (Jan 2024-Present)
- Analyzed large language models’ data curation practices on the AllenNLP team, in collaboration with Luca Soldaini, Emma Strubell, Suchin Gururangan, and Lauren F. Klein.
- Launched a new AI & Education research initiative, funded by the Gates Foundation.

Allen Institute for Artificial Intelligence

Research Intern

- Mentors: Katie Keith, Jesse Dodge
- Mapped scientific domains on the Semantic Scholar and AllenNLP teams.
- Awarded “Outstanding Intern of the Year,” and published our paper in Findings of ACL 2023.

Microsoft Research

Research Intern

- Mentors: Alexandra Olteanu, Su Lin Blodgett
- Evaluated natural language generation systems on the Fairness, Accountability, Transparency, and Ethics (FATE) team, in collaboration with Milad Shokouhi and Hanna Wallach.
- Presented our paper at the Text as Data conference.

Stanford Computer Science

Research Assistant

- Advisors: Dan Jurafsky, Patricia Bromley.
- Investigated the framing and representation of underrepresented groups in history textbooks with linguistics PhD student Dora Demszky.
- Won a best paper award.

École Polytechnique Fédérale de Lausanne

Research Intern

- Advisor: Robert West (Data Science Lab)
- Operationalized and analyzed behavioral trends in a political quote dataset using Apache Spark, emotion lexicons, Stanford CoreNLP parsers, and social networks.

Stanford Computer Science

Research Assistant

- Advisors: David Jurgens, Jure Leskovec (Stanford Network Analysis Project), Dan Jurafsky (NLP group)
- Leveraged language and social network features to classify fictional and real relationships with scikit-learn, NLTK, and Keras.

Teaching and Mentoring Experience

Stanford CS 224U, Natural Language Understanding, Course Assistant (top 5% in CS)

Symbolic Systems Program, Advising Fellow

Stanford EE/CME 103, Introduction to Matrix Methods, Course Assistant

Undergraduate advisees: Odelia Larbi-Amoah (Emory, Current), Sabrina Baur (UC Berkeley, Current), Ethan N. Elasky (UC Berkeley, Current), Claire Wang (UC Berkeley, 2022-Current), JJ Kim-Ebio (UC Berkeley, 2022-Present), Tiffany Liang (UC Berkeley, 2023), Aryia Dattamajumdar (UC Berkeley, 2023), Sebastian Orozco (UC Berkeley, 2022), Divya Tadimeti (UC Berkeley, 2021-2022), Nikhil Mandava (UC Berkeley, 2021).

Master’s student advisees: Vyoma Raman (Stanford, Current).

Service

Professional

Reviewer: ACL Rolling Review (2021-Present), ACL (2021-Present), EMNLP (2021-Present), FAccT (Present), Journal of Big Data (2024), LREC (2024), NeurIPS (2023), AI & HCI Workshop at ICML (2023), The Web Conference (2023), NLP for Positive Impact Workshop (2022), COLING (2022), NAACL Student Research Workshop (2022), NAACL Workshop on Understanding Implicit and Underspecified Language (2022), SCIL (2022), CHI (2022), CSCW (2022), ACL Workshop on NLP for Positive Impact (2021), AERA Open (2020), NeurIPS Human and Machine in-the-Loop Evaluation and Learning Strategies Workshop (2020).

Organizing committee: Teaching NLP (2021) at NAACL.

Advisory board: AI for Humanists (2020-Present), UC Berkeley Human Rights Center's Assessment of Large Language Models (2023-Present).

Community

General: Diaries of Social Data Research (Podcast Host, 2021-Present), Sociologists of Digital Things (Admin, 2021), NLP+CSS PhD Summer Reading Group (2020).

Students: Berkeley Undergraduate Research Apprentice Program (2021-Present), BAIR Mentoring Program (2020-2021), CS Kickstart (Speaker; 2020), UC Berkeley Girls in Engineering (Leader; 2020), Berkeley AI4ALL (Mentor; 2019), Stanford AI4All (Mentor; 2019), Girls Teaching Girls to Code (Mentor/Lead; 2018, 2019).

Skills

Computer Languages: Python, Julia, SQL

Natural Languages: English, Mandarin Chinese

Tools: NLTK, Stanford CoreNLP, SpaCy, scikit-learn, Apache Spark, MTurk, Figure Eight, Keras, TensorFlow, PyTorch.